

# CS470: Introduction

Instructor: Seunghoon Hong

# Agenda

- Course logistics
- Introduction to AI
- Course preview

# Teaching staffs

## Instructor:

Seunghoon Hong [Room] 3429, E3-1 [email] [seunghoon.hong@kaist.ac.kr](mailto:seunghoon.hong@kaist.ac.kr)

## TAs:

Whie Jung [email] [whieya@kaist.ac.kr](mailto:whieya@kaist.ac.kr)

Sunghyun Myung [email] [shm\\_@kaist.ac.kr](mailto:shm_@kaist.ac.kr)

Hajin Shim [email] [shimazing@kaist.ac.kr](mailto:shimazing@kaist.ac.kr)

# Course logistics

- Course materials
  - Course website: KLMS
    - Slides, announcements, assignments, project.
  - Optional textbook: [Goodfellow et al., Deep learning, MIT Press, 2016](#)

# Course logistics

- Prerequisites
  - Proficiency in python programming and linear algebra

# Course logistics

- Grading policy
  - Attendance (5%)
  - Quiz (10%)
  - Midterm (25%)
  - Programming assignments (30%)
  - Final project (30%)

# Assignments

- Programming platform:
  - Google Colab + Pytorch
  - Other libraries such as Tensorflow are also allowed, but not recommended since most course materials will be prepared based on Pytorch
- Late submission
  - Submission deadline: **midnight** at every due date
  - Late submissions are allowed until the **three days** after the on-time deadline with **20% penalty**.
  - You have total two days of grace period.

# Final project

Form a team up to 4 people.

Define and solve your own problem



# Final project

Define and solve your own problem.

- Solve the **novel task** by applying the existing models  
(e.g. developing a web-service, interactive demo, mobile application, game, benchmark, etc.)
- You can use the existing implementations or even the pre-trained models, but your **task** (and hopefully evaluation) should be novel and well-defined.
- Detailed grading criteria will be announced shortly, but we encourage the followings:
  - Your task is novel and/or useful and/or fun (yes, it is subjective but will be very generous 😊).
  - The evaluation (or analysis) on your task is comprehensive and convincing.
  - Your approach has technical contributions  
(e.g. non-trivial modifications to the existing model, your own implementation, etc.)

# Final project

## Deliverables

- **Code:** a link to the Github repository
- **Project report:** 4-page report (in english)

## Schedule

- Team setup: 9/11
- Project proposal: 10/16
- Project presentation: 11/30 - 12/9 (4 classes)
- Code and report submission deadline: 12/13

# Academic Integrity

- All individual assignment should be completed entirely on your own.  
Discussions between classmates are allowed and encouraged,  
but sharing or seeing other students' solutions is not allowed.
- For team project, you should specify the contributions of each person.  
Contributions should be clear and specific.  
Example: X implemented these functions, Y conducted evaluation on dataset A and B,  
Z performed hyper-parameter tuning on C, D, E, etc.
- Failure to adhere to these policies may lead to serious penalties (e.g. course failure)

# Office hour

- Main communication channel: **KLMS Q&A board**
- **TAs:** Wed 6:30 pm - 9:30 pm (via google meet. Link will be shared later)
- **Professor:** by appointment

# Introduction to AI

# AI in the wild

- Self-driving cars



## Waymo Can Finally Bring Truly Driverless Cars to California

The company born as Google's self-driving car project is the first with the right to test human-free cars on public roads in the Golden State.



# AI in the wild

- Superhuman performance in games



IBM Watson (2011, Jeopardy!)



Deepmind AlphaGo (2016, Go)

Not yet but getting closer on ...




OpenAI Five (2018, League of Legends)



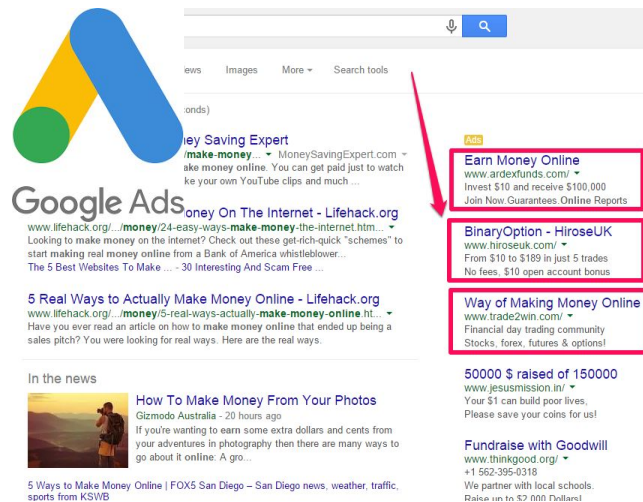
Deepmind AlphaStar (2018, StarCraft 2)

# AI in the wild

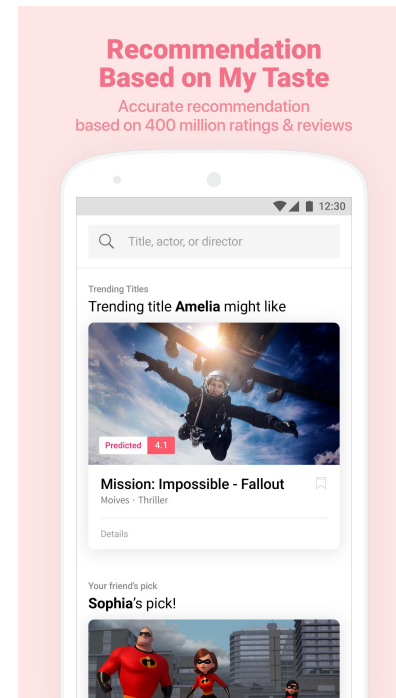
- Recommendations



YouTube



Google Ad

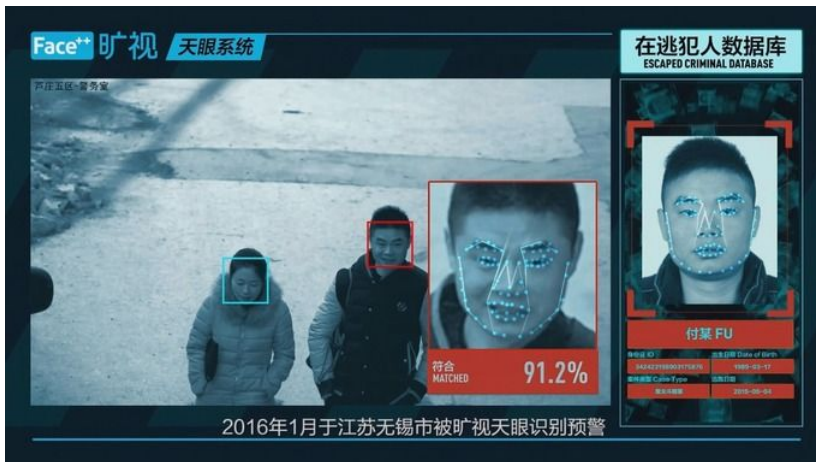


Watcha




# AI in the wild

- Security



# AI in the wild

- Healthcare




IBM Watson Health

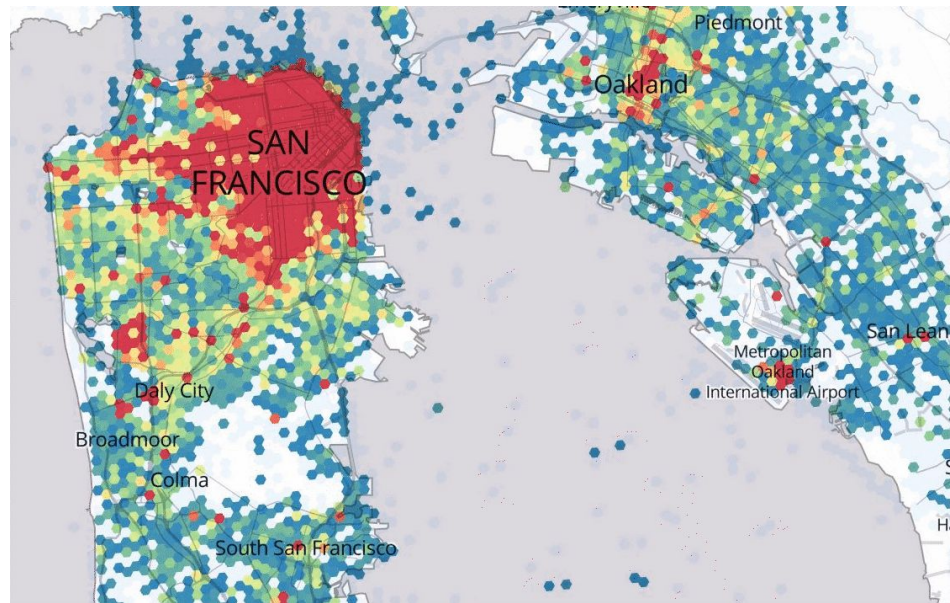


# AI in the wild

- Economy




Reducing the energy usage of Google data centers (Deepmind)



Predicting supply/demand in Uber

# AI in the wild

- Daily assistance



Virtual assistant  
(Apple siri, Google assistance)



Spell checker (grammarly)

# There are speculations/hypes on AI

The Atlantic SUBSCRIBE SEARCH MENU

BUSINESS

Technology

## Stephen Hawking warns artificial intelligence could end mankind

By Rory Cellan-Jones  
Technology correspondent

2 December 2014 | Technology

Elon Musk  
It  
"It's definitely  
BY TESS T  
f t g+

ERIK BRYNJOLFS  
ANDREW MCAFEE

Stephen Hawking: "Humans, who are limited by slow biological evolution, couldn't compete and would be superseded"

IMAGE: Getty Images

THE ADVANCES WE'VE MADE IN AI, FROM HUMANOID ROBOTS, SPEECH RECOGNITION AND SYMBIOTICALLY-CONTROLLED ROBOTS TO *Jeopardy!*-CHAMPION COMPUTERS—ARE NOT THE

Elon Musk has emerged as a leading voice in speaking out on the potential dangers of artificial intelligence, going so far as to call it the "biggest existential threat" to humanity.

HIT THE ROAD

GEHRY-DESIGNED NEW YORK CITY OFFICE

# The objective of this course

- Provide an overview of recent interesting progresses in AI
- Introduce some principles/techniques that has enabled these advances
- Help you understand frontiers/limitations in the current technology
- Hopefully encourage you to find/solve your own problems

So, what is Artificial Intelligence?

# What is an intelligence?

- An ability to perceive, reason, act, and learn.



# What is an intelligence?

- An ability to perceive, reason, act, and learn.




# What is an intelligence?

- An ability to **perceive**, reason, act, and learn.




Observe the scene through  
your visual cortex



# What is an intelligence?

- An ability to perceive, **reason**, act, and learn.



There are people crossing the road

This woman may walk slowly due to her companions

Maybe we should not enter this area

speed limit is 30mph

# What is an intelligence?

- An ability to perceive, reason, **act**, and learn.



Let's make a left turn!

# What is an intelligence?

- An ability to perceive, reason, act, and **learn**.



Ok. It was a wrong direction.  
I should have turned right.

# The goal of AI

- Modeling the *components of intelligence* using computable functions

# The goal of AI

- Modeling the *components of intelligence* using computable functions

## Perception

convert the observation into a  
machine-understandable form



$x$

# The goal of AI

- Modeling the *components of intelligence* using computable functions

## Perception

convert the observation into a machine-understandable form



$x$

## Action

define decisions to achieve the task

Turn right  
Turn left  
Go forward  
Go backward  
Stop

$y$



# The goal of AI

- Modeling the *components of intelligence* using computable functions

## Perception

convert the observation into a machine-understandable form



$x$

## Reason

design a mapping between data and output

$f(x)$

## Action

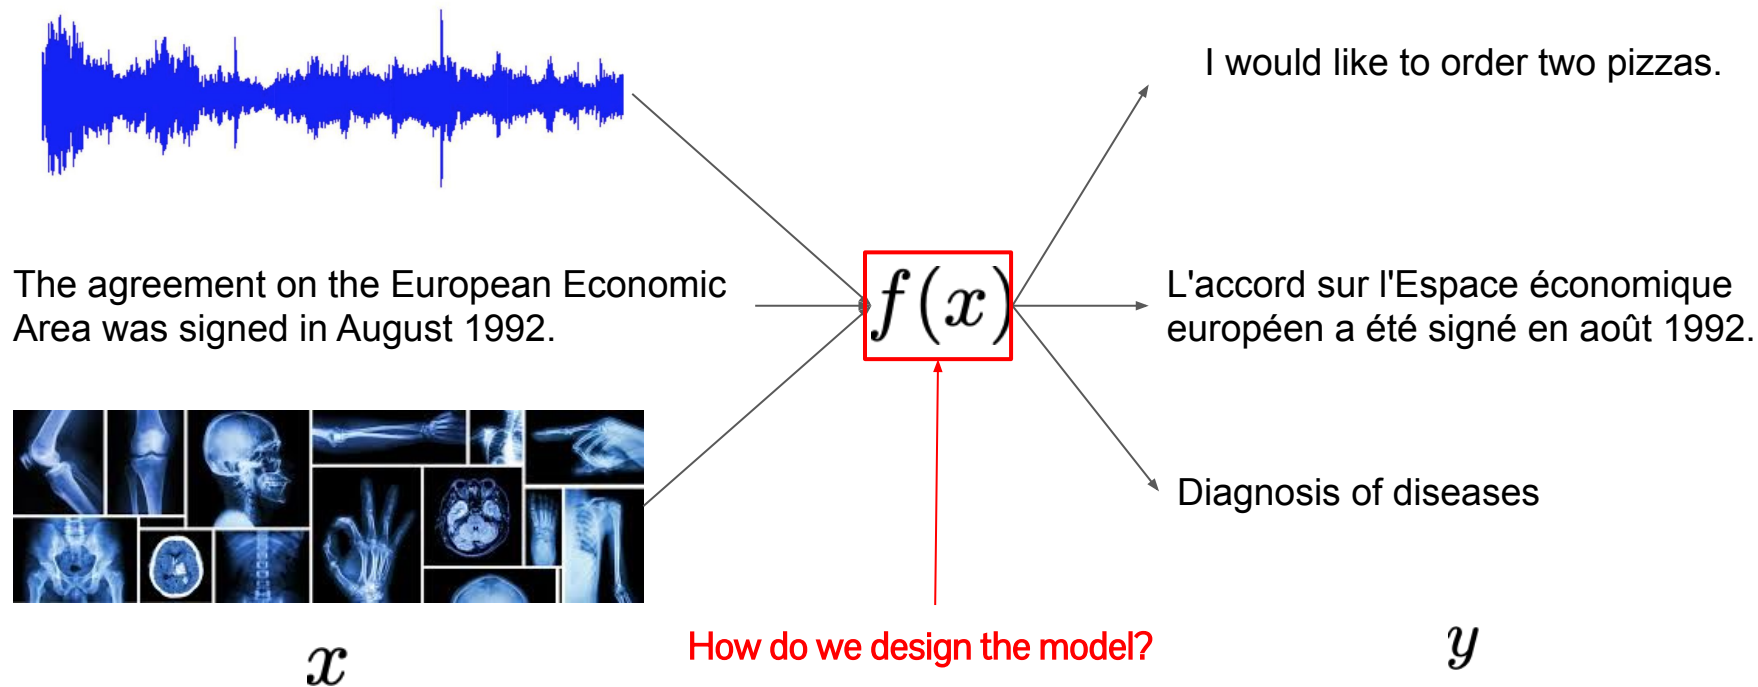
define decisions to achieve the task

Turn right  
Turn left  
Go forward  
Go backward  
Stop

$y$

# The goal of AI

- Modeling the *components of intelligence* using computable functions



# AI and ML

## ARTIFICIAL INTELLIGENCE

Any technique that enables  
computers to mimic  
human behavior



# AI and ML

## ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior



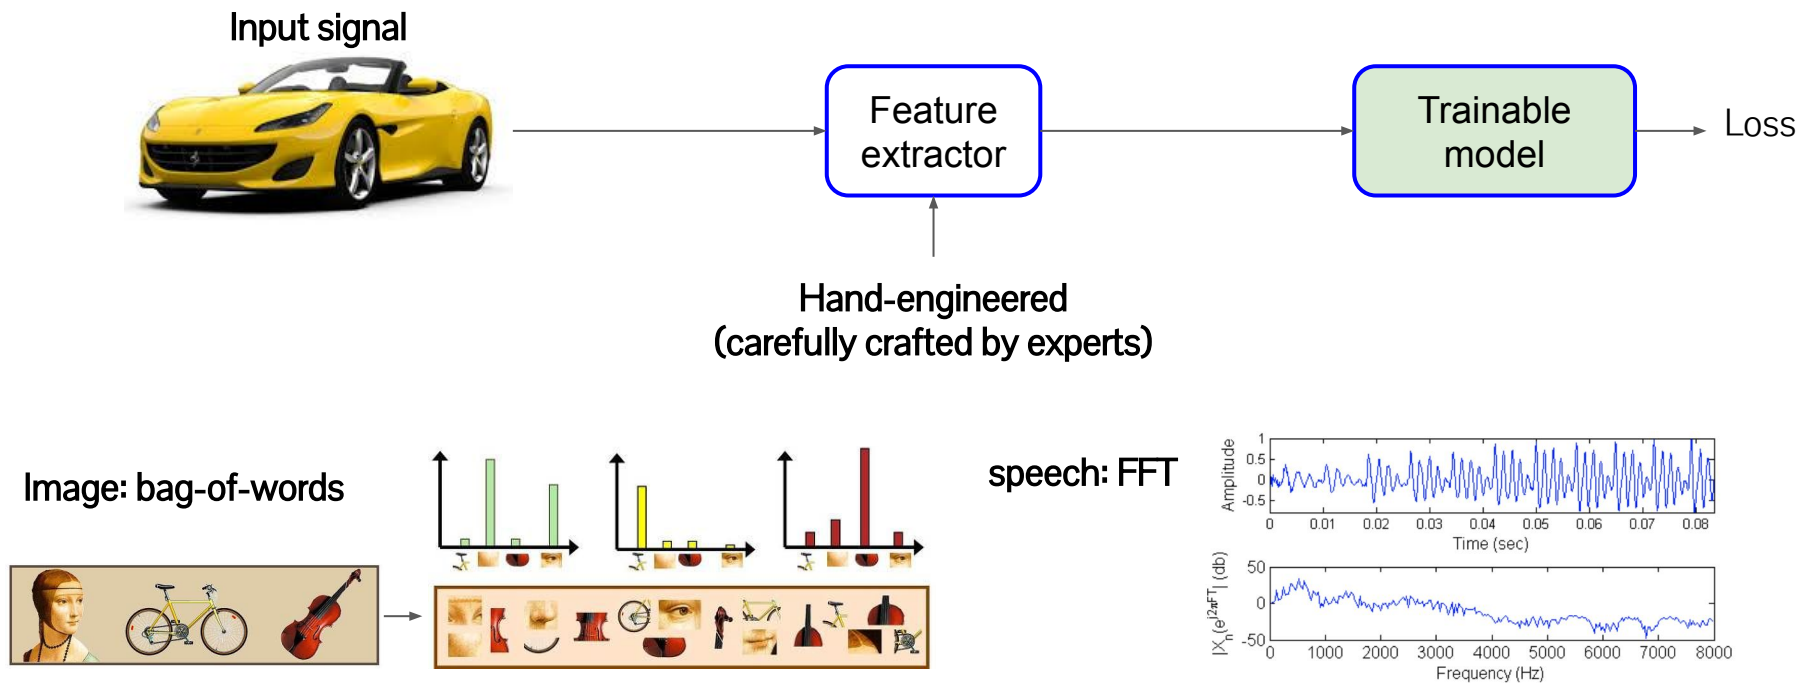
## MACHINE LEARNING

Ability to learn without explicitly being programmed




# ML workflow

- Traditional pattern recognition

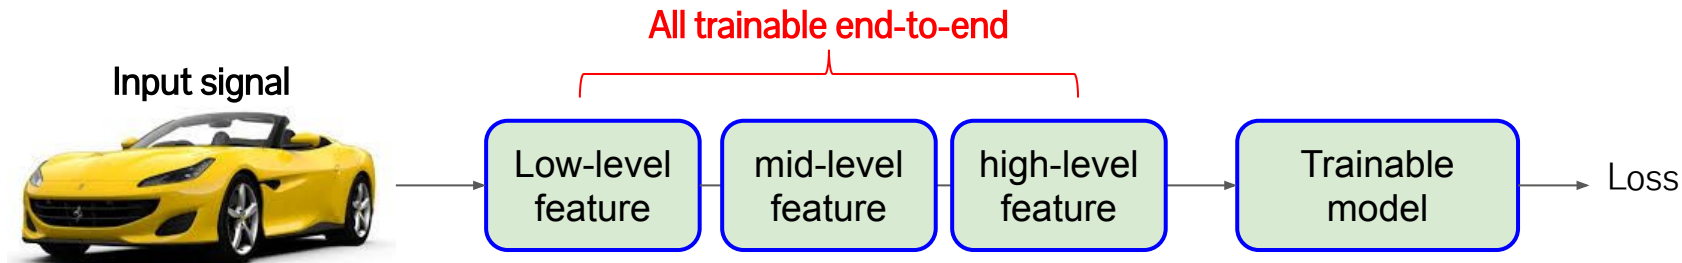


# ML workflow


- Traditional pattern recognition



- Deep learning




# Preview: deep learning representations



# DL revolutionized ML


Top 5 classification error on ImageNet






# DL revolutionized ML

- Speech recognition over time



# DL revolutionized ML

## Growing Use of Deep Learning at Google



# Course overview



# Course overview

- ❑ Image classification
- ❑ Object detection
- ❑ Semantic segmentation
- ❑ Pose estimation
- ❑ Visualization
- ❑ Style transfer
- ❑ Adversarial attacks
- ❑ Text classification
- ❑ Sentiment analysis
- ❑ Machine translation
- ❑ Sketch synthesis
- ❑ Image captioning
- ❑ Visual question answering
- ❑ Image generation
- ❑ Text generation
- ❑ Text-to-image synthesis
- ❑ Img-to-img translation
- ❑ Unpaired img-to-img translation
- ❑ Interactive drawing
- ❑ Search algorithms
- ❑ Markov decision processes
- ❑ Reinforcement learning



Preview of the selected applications

# Visual understanding


- Image classification



Person,  
Elephant,  
Sky,  
Field,  
...

# Visual understanding

- Object detection



# Visual understanding

- Semantic segmentation





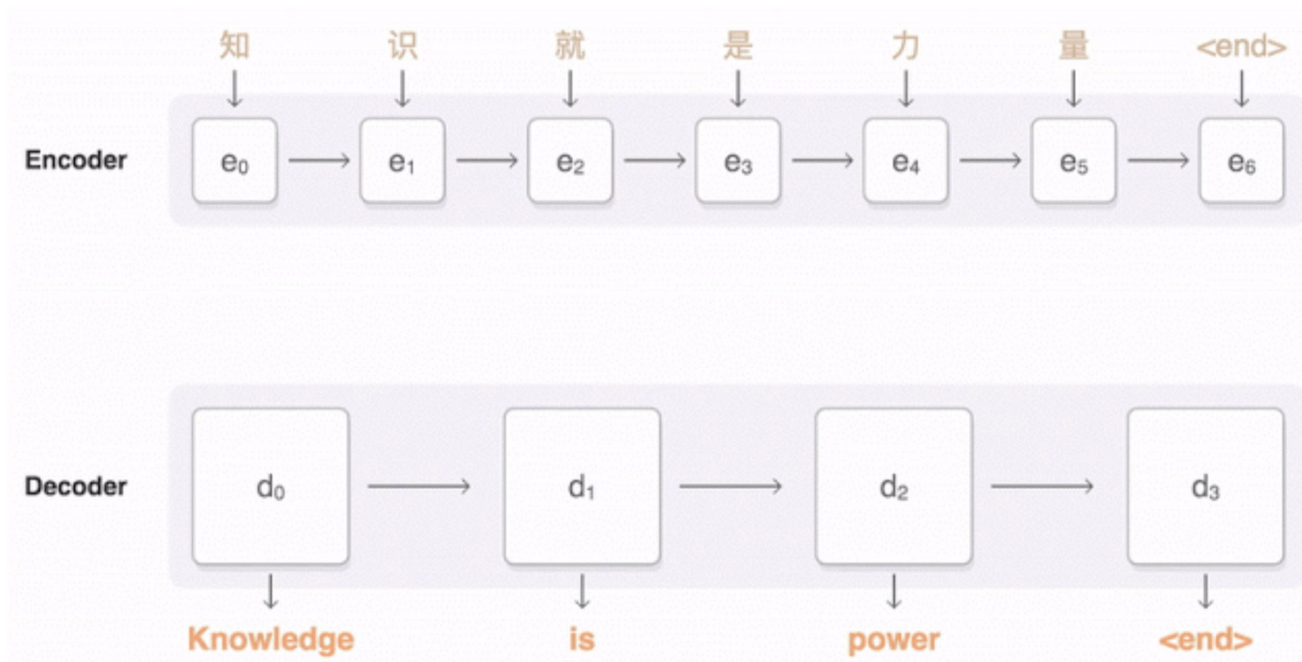
# Visual understanding

- Pose estimation




# Language understanding

- Machine translation



# Audio understanding

- Speech recognition



# Understanding neural network


- Adversarial attack



**"panda"**

57.7% confidence

+  $\epsilon$



=





**"gibbon"**

99.3% confidence



# Understanding neural network

- Style transfer



# Generative modeling

- Generating images



# Generative modeling

- Generating text

SYSTEM PROMPT  
(HUMAN-WRITTEN)

*Miley Cyrus was caught shoplifting from Abercrombie and Fitch on Hollywood Boulevard today.*

MODEL  
COMPLETION  
(MACHINE-  
WRITTEN, SECOND  
TRY)

The 19-year-old singer was caught on camera being escorted out of the store by security guards.

The singer was wearing a black hoodie with the label 'Blurred Lines' on the front and 'Fashion Police' on the back.

Scroll down for video

Shoplifting: Miley Cyrus was caught shoplifting from Abercrombie and Fitch on Hollywood Boulevard today (pictured)

The singer was wearing a black hoodie with the label 'Blurred Lines' on the front and 'Fashion Police' on the back

The singer was also wearing a pair of black-rimmed glasses, a black jacket, black jeans and black sandals.

She was carrying a pair of black and white striped gloves and a small black bag.

# Generative modeling

- Generating audio signals


<https://openai.com/blog/musenet/>






# Conditional generative modeling

- Interactive generation



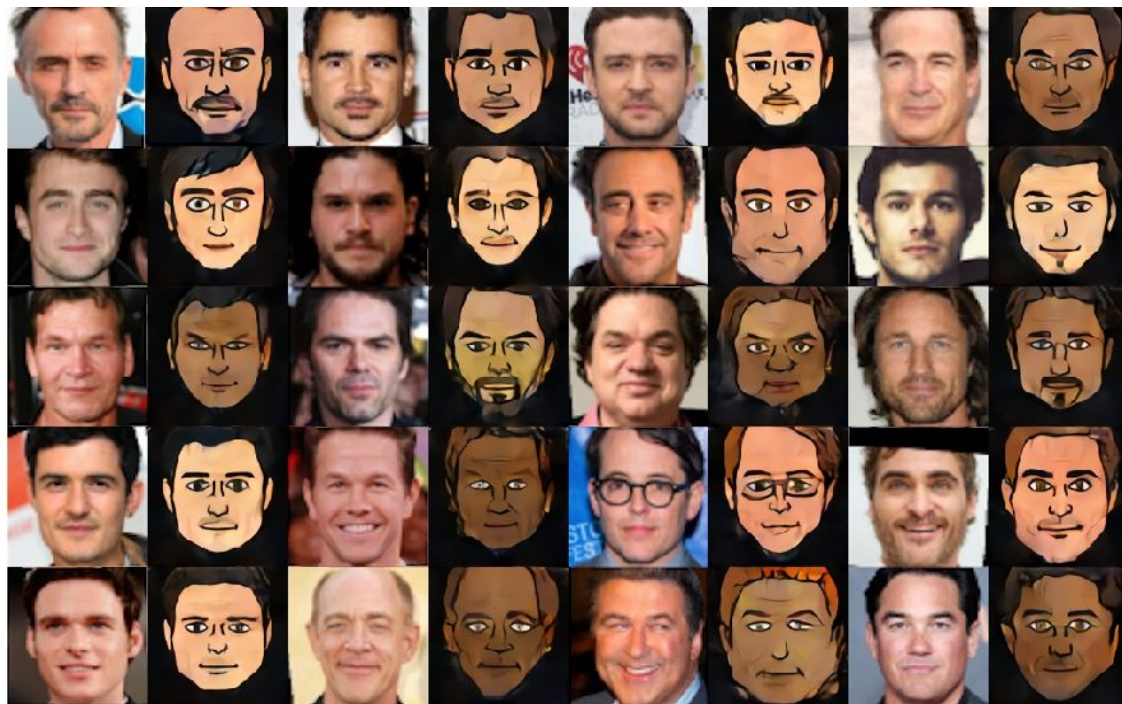
# Unsupervised conditional generative modeling

- Motion transfer



# Unsupervised conditional generative modeling

- Image translation from unpaired data



# Next

- Neural network basics

Any questions?