

CS470 Lab




Tutorials for Colab and Pytorch

Announcement

- Team formation deadline is due this Friday!
- The first assignment will be release in the next week

Lower loss = better model?


- So far, it seems like the neural network training is all about minimizing training error
- But does the lower training error always mean a better model?




Lower loss = better model?

- So far, it seems like the neural network training is all about minimizing training error
- But does the lower training error always mean a better model?


Training loss (error)



>




>




Lower loss = better model?

- So far, it seems like the neural network training is all about minimizing training error
- But does the lower training error always mean a better model?


Training loss (error)



>



>



There are also arbitrary many solutions that achieves the similar loss: **which one is better?**

Lower loss \neq better model

- Overfitting (memorization):
 - The model simply “memorizes” the training examples
 - Achieves very low training error, but very high test error (not generalized to unseen examples)
 - This problem is prevalent especially when **# of parameters \gg # of training data**
(can fit arbitrary complex functions to the data)

Regularizing neural network

- Option 1: weight decay
- Option 2: dropout
- Option 3: early stopping

Weight decay

- Penalize complex solutions using additional constraints

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \frac{1}{N} \sum_{I=1}^N \mathcal{L}(f(\mathbf{x}^{(i)}; \mathbf{W}), \mathbf{y}^{(i)}) + \lambda R(\mathbf{W})$$

Training loss: how well the model fit to the training data

Regularization: add constraints to make the model behaves well

Weighting parameter: determine Importance of regularization

Weight decay

- Penalize complex solutions using additional constraints

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \frac{1}{N} \sum_{I=1}^N \mathcal{L}(f(\mathbf{x}^{(i)}; \mathbf{W}), \mathbf{y}^{(i)}) + \lambda R(\mathbf{W})$$

Examples of regularization


- **L2** regularization $R(\mathbf{W}) = \|\mathbf{W}\|_2$: Prefer the weights roughly spreaded over all neurons (i.e. make use of all neurons equally important)
- **L1** regularization $R(\mathbf{W}) = \|\mathbf{W}\|_1$: Prefer sparse weights (i.e. less complicated functions)

Weight decay

- Penalize complex solutions using additional constraints

$$\mathbf{W}^* = \arg \min_{\mathbf{W}} \frac{1}{N} \sum_{I=1}^N \mathcal{L}(f(\mathbf{x}^{(i)}; \mathbf{W}), \mathbf{y}^{(i)}) + \lambda R(\mathbf{W})$$


Training loss (error)



>


<

Training loss + λ regularization




Dropout

- Randomly **turn off** activations with some probability p




Dropout

- Randomly **turn off** activations with some probability p



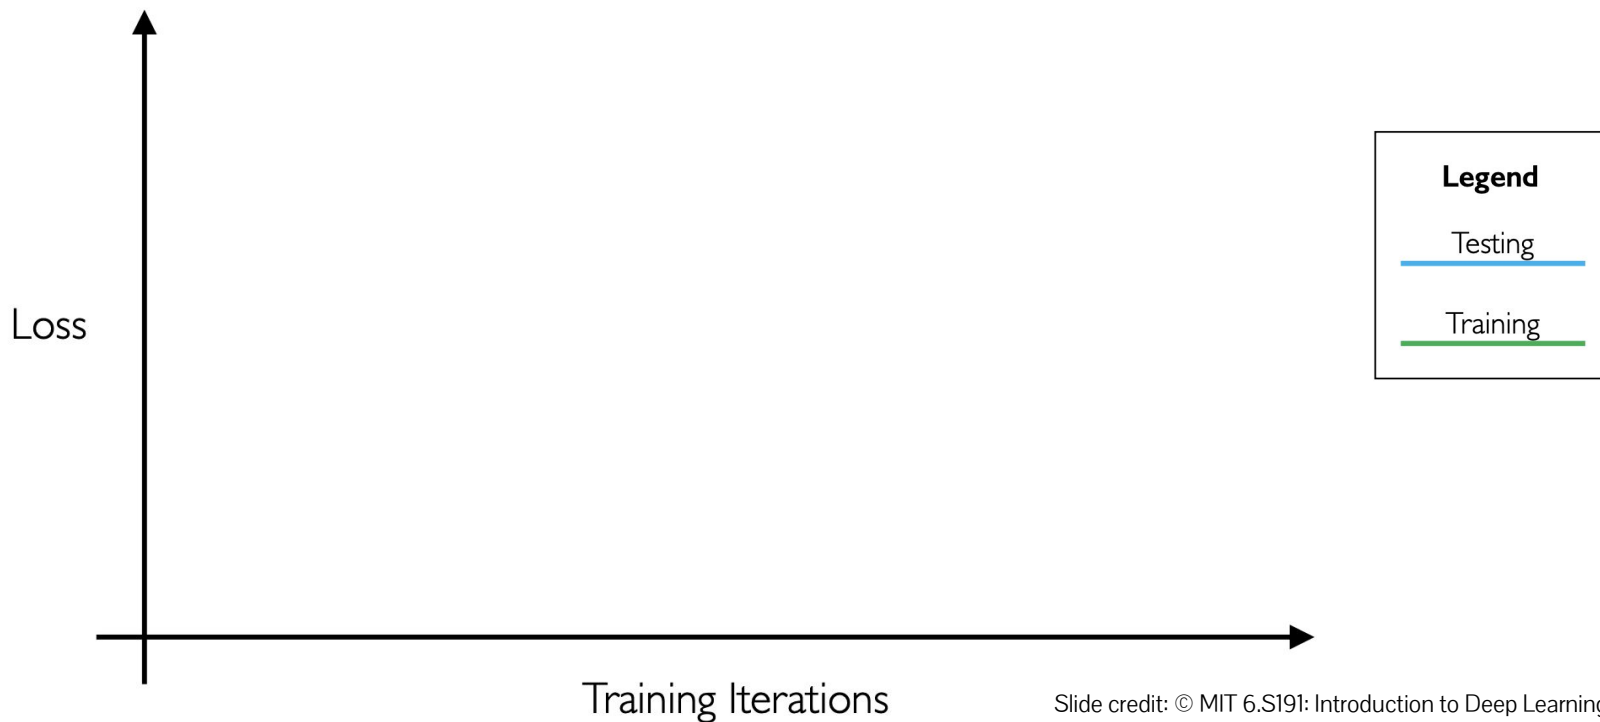
Dropout

- Randomly **turn off** activations with some probability p
 - **Intuition:** add stochasticity to the network to prevent memorization
(every forward propagation leads to different outputs even for the same input)




Early stopping

- Stop training once the validation loss starts to increase



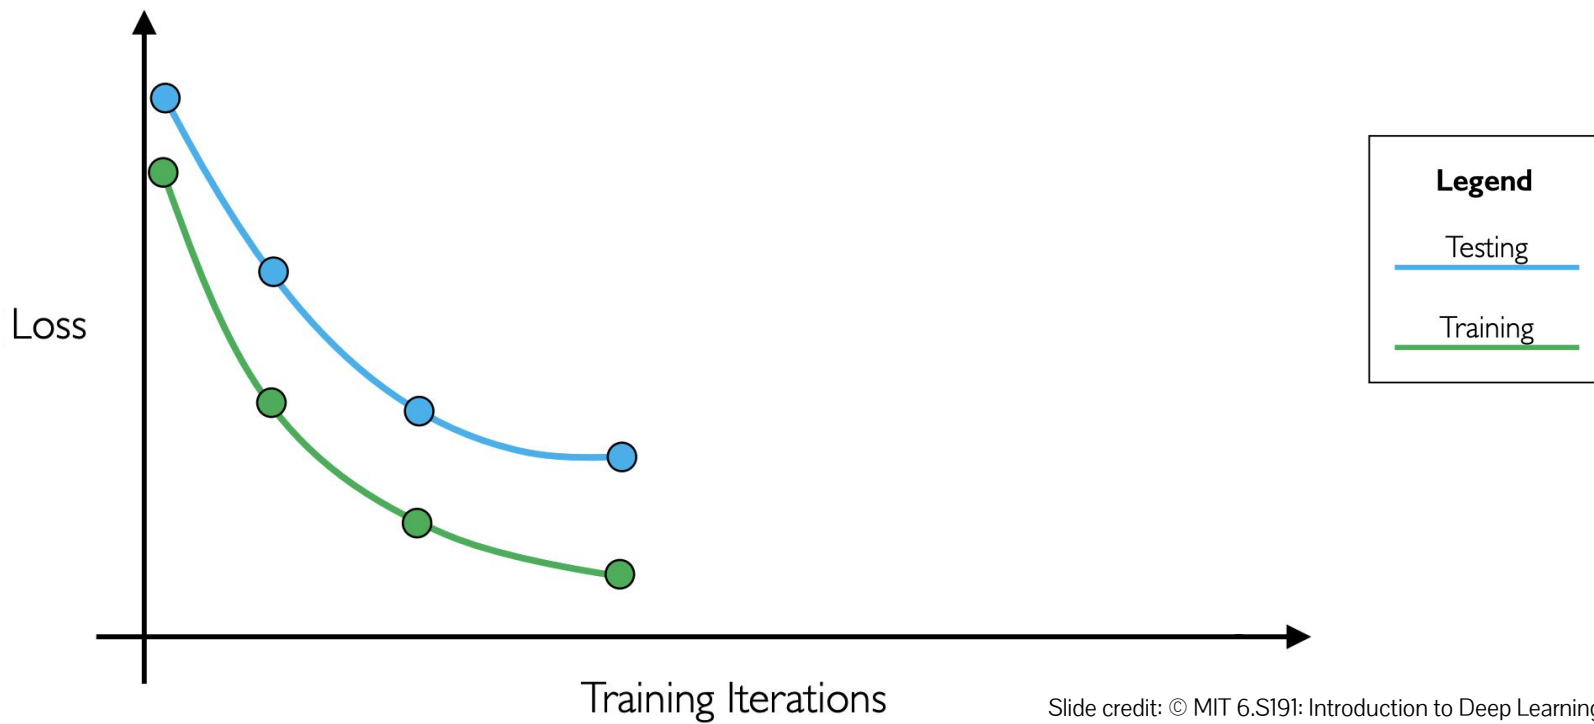
Early stopping

- Stop training once the validation loss starts to increase




Early stopping

- Stop training once the validation loss starts to increase




Early stopping

- Stop training once the validation loss starts to increase




Early stopping

- Stop training once the validation loss starts to increase




Early stopping

- Stop training once the validation loss starts to increase



Early stopping

- Stop training once the validation loss starts to increase



Summary:

- Overfitting: low training error, high testing error
 - Add regularizations to prevent memorization
 - Popular regularizations: weight decay, dropout, early stopping
- Underfitting: high training/testing error
 - Increase the model capacity/learning rate or train longer
- Improving generalization is an active research area
 - We will also discuss some other approaches in later parts of this course

Table of Contents [\[link to the materials\]](#)

- Google Colaboratory
 - a. Create a google account
 - b. **What is Colaboratory?**
 - c. **How to use GPU in Colab**
 - d. **How to connect a Colab notebook with your google drive**
 - e. **Access to your google drive in Colab**
- Pytorch Tutorial
 - a. Tensor and its basic operations
 - b. Autograd and automatic differentiation
 - c. Building neural networks and optimizers
 - d. Data pipeline
 - e. **Train and Test a simple MLP-based MNIST classifier**

PyTorch Tutorial

Basics components and operations

Importing PyTorch

- Colab supports PyTorch by default

```
[1] 1 import torch
    2
    3 torch.__version__
```

↳ '1.4.0'

Tensors: basic computing unit of Pytorch

- Basically, tensors are for representing scalars, vectors, and matrices
- Similar to NumPy's ndarrays, but supports GPU acceleration

Matrix

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \rightarrow$$

Defining a tensor

```
[1] import torch
import numpy as np

[2] # numpy ndarray
np.array([[1, 2], [3, 4]])

↳ array([[1, 2],
        [3, 4]])

[3] # pytorch tensor on cpu
torch.tensor([[1, 2], [3, 4]])

↳ tensor([[1, 2],
         [3, 4]])
```

Defining a tensor on GPU

```
[4] # pytorch tensor on gpu
torch.tensor([[1, 2], [3, 4]], device='cuda')

↳ tensor([[1, 2],
         [3, 4]], device='cuda:0')
```

Tensor <-> ndarray

```
[5] # numpy ndarray to pytorch tensor
np_array = np.array([[1, 2], [3, 4]])
torch.from_numpy(np_array)

↳ tensor([[1, 2],
         [3, 4]])

[6] # pytorch tensor to numpy ndarray
torch_tensor = torch.tensor([[1, 2], [3, 4]])
torch_tensor.numpy()

↳ array([[1, 2],
        [3, 4]])
```

Basic arithmetic operations with tensors

Addition $\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} + \begin{bmatrix} 5 & 4 \\ 3 & 2 \end{bmatrix} = \begin{bmatrix} 6 & 6 \\ 6 & 6 \end{bmatrix}$

Subtraction $\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} - \begin{bmatrix} 5 & 4 \\ 3 & 2 \end{bmatrix} = \begin{bmatrix} -4 & -2 \\ 0 & 2 \end{bmatrix}$

element-wise multiplication $\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \otimes \begin{bmatrix} 5 & 4 \\ 3 & 2 \end{bmatrix} = \begin{bmatrix} 5 & 8 \\ 9 & 8 \end{bmatrix}$

element-wise division $\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \oslash \begin{bmatrix} 5 & 4 \\ 3 & 2 \end{bmatrix} = \begin{bmatrix} 0.2 & 0.5 \\ 1.0 & 2.0 \end{bmatrix}$

```
[7] a = torch.tensor([[1., 2.], [3., 4.]])  
    b = torch.tensor([[5., 4.], [3., 2.]])
```

```
[8] a+b
```

```
↳ tensor([[6., 6.],  
          [6., 6.]])
```

```
[9] a-b
```

```
↳ tensor([[ -4., -2.],  
          [  0.,  2.]])
```

```
[10] a*b
```

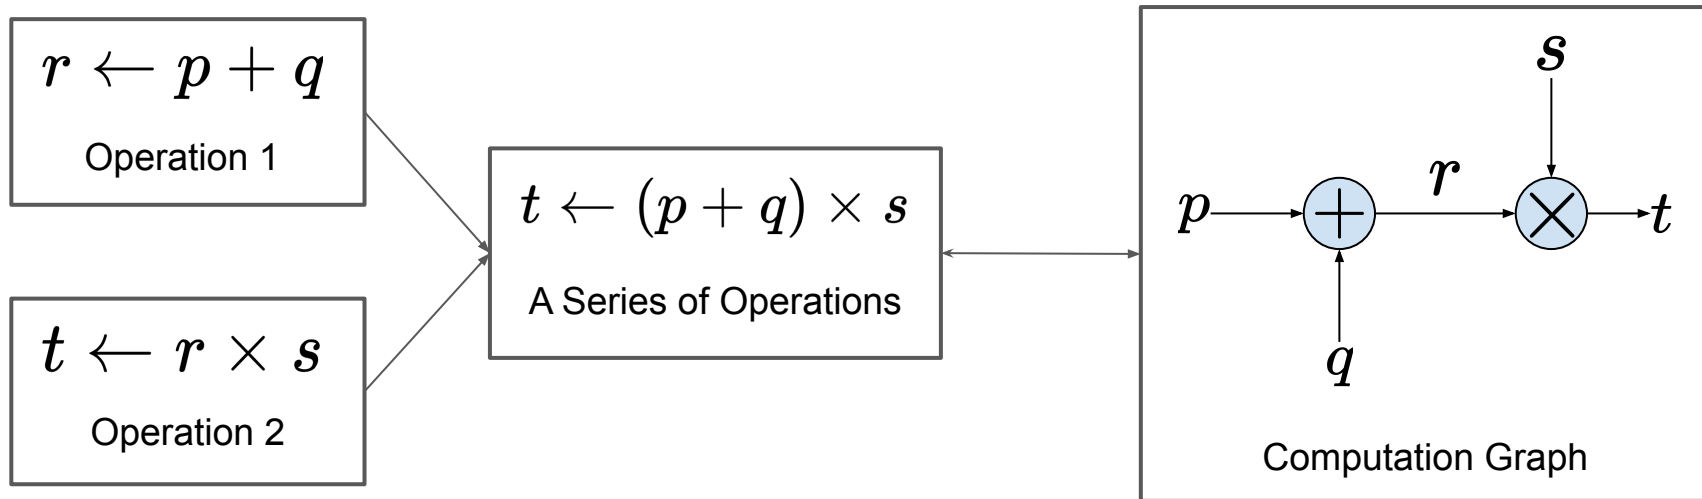
```
↳ tensor([[5., 8.],  
          [9., 8.]])
```

```
[11] a/b
```

```
↳ tensor([[0.2000, 0.5000],  
          [1.0000, 2.0000]])
```


Computation graph

- A series of operations constructs a computation graph
- Any operation between tensors defines a node in the computation graph




Computation graph and Forward function

- Consider below computation graph as our forward function



- If we assume $p=3$, $q=5$, and $s=2$, then we get $r=8$ and $t=16$.




```
[14] p = torch.tensor(3.)  
      q = torch.tensor(5.)  
      s = torch.tensor(2.)  
      r = p + q  
      t = r * s
```


```
print('r : ', r)  
print('t : ', t)
```

```
➡ r : tensor(8.)  
   t : tensor(16.)
```

Forward & Backward functions




Forward Function



Backward Function

Backward function and Chain rule



$$\frac{\partial t}{\partial s} = \frac{\partial(r \times s)}{\partial s} = r \quad \frac{\partial r}{\partial p} = \frac{\partial(p+q)}{\partial p} = 1$$


$$\frac{\partial t}{\partial r} = \frac{\partial(r \times s)}{\partial r} = s \quad \frac{\partial r}{\partial q} = \frac{\partial(p+q)}{\partial q} = 1$$

$$\frac{\partial t}{\partial q} = \frac{\partial t}{\partial r} \times \frac{\partial r}{\partial q} = s \times 1 = s$$

$$\frac{\partial t}{\partial p} = \frac{\partial t}{\partial r} \times \frac{\partial r}{\partial p} = s \times 1 = s$$

Backward function: an example

- If we assume $p=3$, $q=5$, $s=2$, $r=8$, and $t=16$, then $dt/ds=8$, $dt/dp=2$ and $dt/dq=2$



Automatic differentiation (AutoGrad)

- Wait, then do we have to calculate all the derivatives on our own ?
- What if the variables are vectors and matrices, but not scalars ?
- No worries ! AutoGrad Package in PyTorch will do that for us.

```
1 import torch
2
3 # Forward Propagation
4 p = torch.tensor([3.], requires_grad=True)
5 q = torch.tensor([5.], requires_grad=True)
6 s = torch.tensor([2.], requires_grad=True)
7 r = p + q
8 t = r * s
9
10 print('p :', p.item())
11 print('q :', q.item())
12 print('s :', s.item())
13 print('r :', r.item())
14 print('t :', t.item())
```

```
p : 3.0
q : 5.0
s : 2.0
r : 8.0
t : 16.0
```

```
16 # Backward Propagation
17 t.backward()
18
19 print('dt/dp :', p.grad.item())
20 print('dt/dq :', q.grad.item())
21 print('dt/ds :', s.grad.item())
```


```
dt/dp : 2.0
dt/dq : 2.0
dt/ds : 8.0
```


PyTorch Tutorial

Building a neural network

Let's build a one-layer baby network

- One-layer classifier for MNIST digit classification




$$\mathbf{y} = \mathbf{x}\mathbf{W} + \mathbf{b}$$

Size of \mathbf{W} ? [784, 10]

Size of \mathbf{b} ? [1, 10]

Let's build a one-layer baby network

- One-layer classifier for MNIST digit classification




$$y = xW + b$$

```
import math
```

```
weights = torch.randn(784, 10) / math.sqrt(784)  
weights.requires_grad_()  
bias = torch.zeros(10, requires_grad=True)
```

Let's build a one-layer baby network

- One-layer classifier for MNIST digit classification



$$\mathbf{y} = \mathbf{x}\mathbf{W} + \mathbf{b}$$

```
import math
```

```
weights = torch.randn(784, 10) / math.sqrt(784)
weights.requires_grad_()
bias = torch.zeros(10, requires_grad=True)
```

```
def log_softmax(x):
    return x - x.exp().sum(-1).log().unsqueeze(-1)

def model(xb):
    return log_softmax(xb @ weights + bias)
```

Baby network: forward propagation

```
bs = 64  # batch size

xb = x_train[0:bs]  # a mini-batch from x
preds = model(xb)  # predictions
preds[0], preds.shape
print(preds[0], preds.shape)
```

output:

```
tensor([-1.9759, -2.1991, -1.9989, -2.4762, -2.6573, -2.2036, -2.7582, -2.5692,
        -2.2971, -2.2089], grad_fn=<SelectBackward>) torch.Size([64, 10])
```

Baby network: loss function

- Binary cross-entropy loss

```
def nll(input, target):  
    return -input[range(target.shape[0]), target].mean()  
  
loss_func = nll
```

Baby network: a training loop

Sampling the minibatch (of the size bs)

Forward & loss computation

Gradient update step

```
from IPython.core.debugger import import set_trace

lr = 0.5 # learning rate
epochs = 2 # how many epochs to train for

for epoch in range(epochs):
    for i in range((n - 1) // bs + 1):
        # set_trace()
        start_i = i * bs
        end_i = start_i + bs
        xb = x_train[start_i:end_i]
        yb = y_train[start_i:end_i]
        pred = model(xb)
        loss = loss_func(pred, yb)

        loss.backward()
        with torch.no_grad():
            weights -= weights.grad * lr
            bias -= bias.grad * lr
            weights.grad.zero_()
            bias.grad.zero_()
```

What should we have for this single-layer network?

- Network parameters
 - Tensors for weight and bias
- Forward and backward mechanisms
 - $y = xW + b$ in this case
- Gradient of parameters
 - All parameters in the network should hold the gradient of the loss w.r.t itself
- Loss function
 - A binary cross entropy loss in this case
- Optimizations
 - Weight initialization, a naive gradient update mechanism (SGD)

How about more complicated networks?

A single linear network



of layers: 1

A matrix multiplication and addition

Inception (GoogleNet)



of layers: a lot

Parallel convolution with different filter sizes, nonlinear functions, batch norm, average and max poolings, multi-head loss, ...

Solution: modularize the computations

- Modularize a layer using a **class**, which comes with handy utility functions (e.g. managing parameters/gradients, forward/backward, switching b/w training/evaluation modes, etc)

```
from torch import nn

class Mnist_Logistic(nn.Module):
    def __init__(self):
        super().__init__()
        self.weights = nn.Parameter(torch.randn(784, 10) / math.sqrt(784))
        self.bias = nn.Parameter(torch.zeros(10))

    def forward(self, xb):
        return xb @ self.weights + self.bias
```

It inherits other utility functions defined in torch.nn.Module

Solution: modularize the computations

- Modularize a layer using a **class**, which comes with handy utility functions (e.g. managing parameters/gradients, forward/backward, switching b/w training/evaluation modes, etc)

```
class Linear(Module):
    __constants__ = ['in_features', 'out_features']

    def __init__(self, in_features, out_features, bias=True):
        super(Linear, self).__init__()
        self.in_features = in_features
        self.out_features = out_features
        self.weight = Parameter(torch.Tensor(out_features, in_features))
        if bias:
            self.bias = Parameter(torch.Tensor(out_features))
        else:
            self.register_parameter('bias', None)
        self.reset_parameters()

    def reset_parameters(self):
        init.kaiming_uniform_(self.weight, a=math.sqrt(5))
        if self.bias is not None:
            fan_in, _ = init._calculate_fan_in_and_fan_out(self.weight)
            bound = 1 / math.sqrt(fan_in)
            init.uniform_(self.bias, -bound, bound)

    def forward(self, input):
        return F.linear(input, self.weight, self.bias)
```

Example:
Actual definition of
fully-connected layer in PyTorch

Solution: modularize the computations

- Modularize a layer using a **class**, which comes with handy utility functions (e.g. managing parameters/gradients, forward/backward, switching b/w training/evaluation modes, etc)
- We can build a complicated neural network by simply composing these layers

Google Colaboratory

Before we start ...

1. Create your google account if you don't have one
2. Go to this [Link](#) and open [Introduction to Colab.ipynb](#)
3. Click `File` tab -> Click `Save a copy in drive` button
 - This will save the notebook file in your google drive, which is necessary to follow the tutorial
 - Check `Colab Notebooks` directory in your [google drive](#) if the notebook is saved successfully

What is Colaboratory?

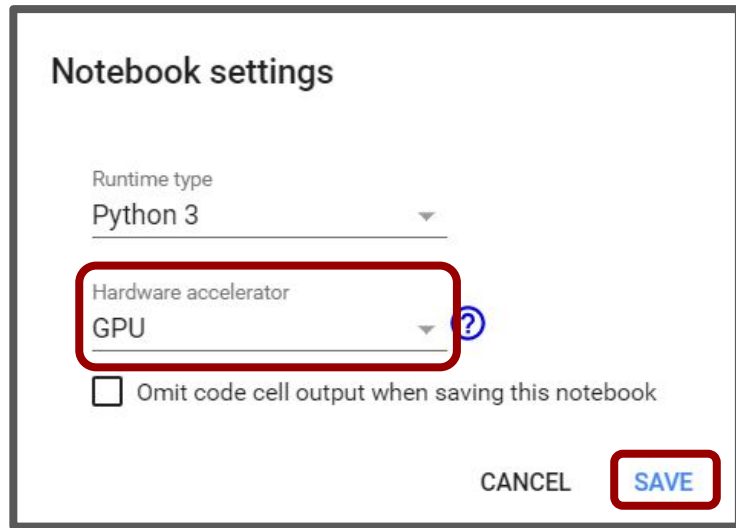

- Colaboratory is a **free Jupyter notebook** environment that requires no setup and runs entirely in the cloud.
- With Colaboratory you can write and execute code, save and share your analyses, and access powerful computing resources, all for free from your browser.
- Colaboratory is run on a Ubuntu 18.04 virtual machine equipped with 13GB RAM, ~310GB Storage limits, and GPUs (K80, TPU).

Handy shortcuts

- Ctrl + M + H : keyboard preferences manager
- Ctrl + M + M : convert to text cell
- Ctrl + M + Y : convert to code cell
- Ctrl + M + A : Insert code cell above
- Ctrl + M + B : Insert code cell below
- Ctrl + M + D : Delete cell/selection
- Shift + Enter : Run cell and select next cell
- Alt + Enter : Run cell and insert new cell
- Ctrl + M + I : Interrupt execution
- Ctrl + M + . : Restart Runtime
- Ctrl + / : comment/uncomment

How to setup GPU

- Runtime -> Change runtime type -> Set Hardware accelerator to GPU -> Save



Mount your google drive - 1

Run the script below, and follow the instruction.

Namely,

1. Go to the given URL in a browser
2. Select your cs470 account and log-in
3. Allow access to the google account
4. Copy the given authorization code, and paste it into the blank below

If you succeed, then you'll see "Mounted at /gdrive"

Note : This step should be repeated everytime you initialize the runtime session


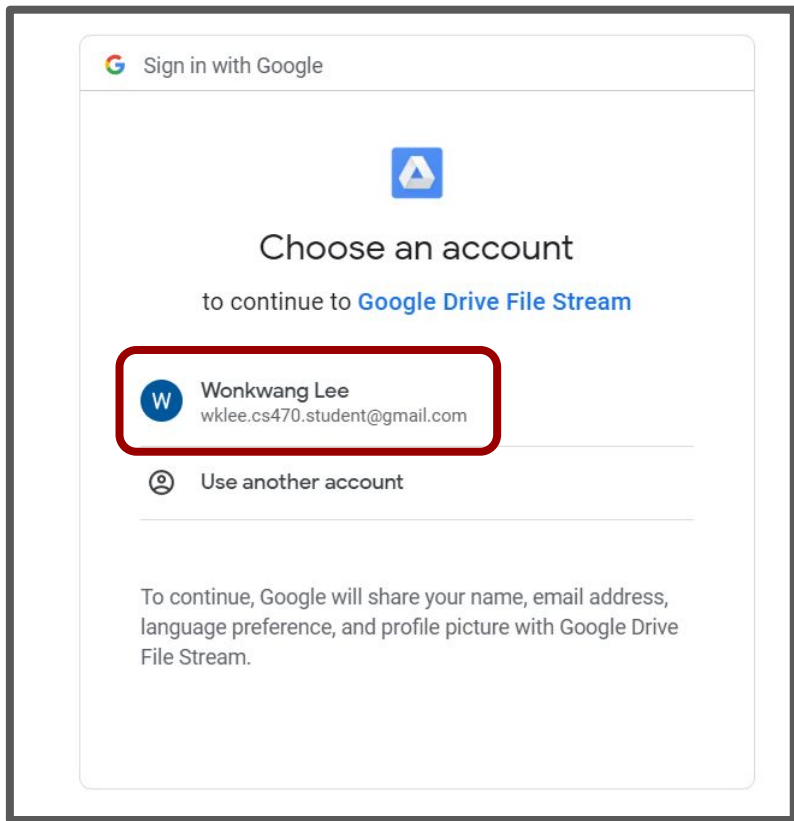


```
from google.colab import drive  
drive.mount('/gdrive')
```

... Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n

Enter your authorization code:

Mount your google drive - 2



Mount your google drive - 3

Run the script below, and follow the instruction.

Namely,

1. Go to the given URL in a browser
2. Select your cs470 account and log-in
3. Allow access to the google account
4. Copy the given authorization code, and paste it into the blank below

If you succeed, then you'll see "Mounted at /gdrive"

Note : This step should be repeated everytime you initialize the runtime sesison



```
from google.colab import drive  
drive.mount('/gdrive')
```

... Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n

Enter your authorization code:

Mount your google drive - 4

```
from google.colab import drive  
drive.mount('/gdrive')
```

... Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6

Enter your authorization code:

.....

```
from google.colab import drive  
drive.mount('/gdrive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.g

Enter your authorization code:

.....

Mounted at /gdrive

Access to your google drive in Colab



```
# check what's in the mounted gdrive using Colab
import os
```

```
gdrive_root = '/gdrive/My Drive'
print('In gdrive :', os.listdir(gdrive_root))
```

```
notebook_dir = os.path.join(gdrive_root, 'Colab Notebooks')
print('In Colab Notebooks :', os.listdir(notebook_dir))
```



```
In gdrive : ['Colab Notebooks']
In Colab Notebooks : ['Copy of Introduction to Colab.ipynb']
```

Download an image into your google drive

```
# download and save an image
!wget https://cs.kaist.ac.kr/common/images/header/logo_top.png -O '/gdrive/My Drive/cs_kaist.png'
print('In gdrive :', os.listdir(gdrive_root))

# Go to the google drive homepage(https://drive.google.com/drive/my-drive),
# log-in using your CS470 account,
# and browse your gdrive directory to check if the image is downloaded successfully

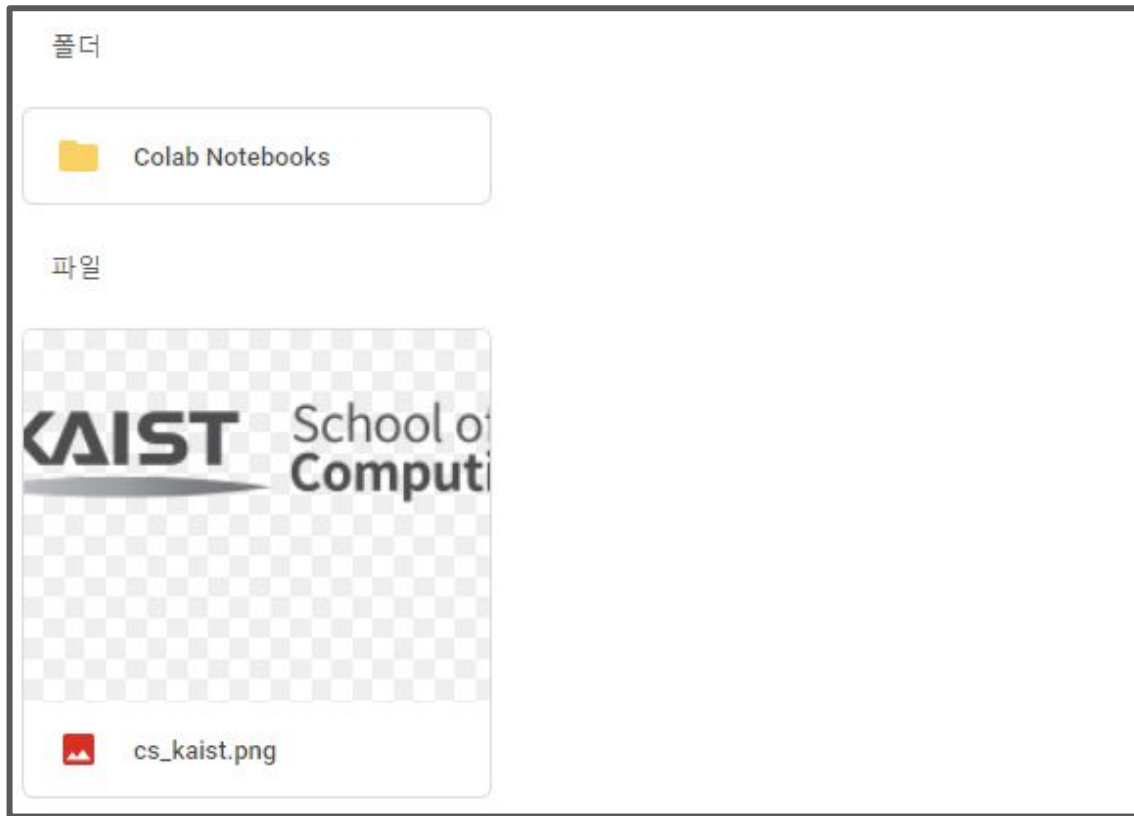
--2019-09-08 12:12:39-- https://cs.kaist.ac.kr/common/images/header/logo\_top.png
Resolving cs.kaist.ac.kr (cs.kaist.ac.kr)... 192.249.19.36
Connecting to cs.kaist.ac.kr (cs.kaist.ac.kr)|192.249.19.36|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 6313 (6.2K) [image/png]
Saving to: '/gdrive/My Drive/cs_kaist.png'

/gdrive/My Drive/cs 100%[=====>]    6.17K  --.-KB/s    in 0s

2019-09-08 12:12:41 (13.3 MB/s) - '/gdrive/My Drive/cs_kaist.png' saved [6313/6313]

In gdrive : ['Colab Notebooks', 'cs_kaist.png']
```

Download an image into your google drive



Load the image from your google drive

```
# load the saved image
from PIL import Image

image_path = os.path.join(gdrive_root, 'cs_kaist.png')
img = Image.open(image_path)
img
```



Disclaimer

- Runtime session will last **at most** 12 hours, regardless of devices (e.g. CPU, GPU, and TPU)
 - if you left the browser opened, probably the session will last at most 12 hours
 - if you closed the browser, probably the session will last at most 90 minutes
- Therefore, it is **highly recommended** that you periodically **back-up your data/outputs to your gdrive** and resume your training by re-loading your saved data. Otherwise, you'll lose everything you've trained as soon as the session is recycled.

PyTorch + Colab

Train and test a simple MLP-based MNIST classifier

Again,

1. Create your google account if you don't have one
2. Go to this [Link](#) and open
[5. Train and Test a simple MLP-based MNIST classifier.ipynb](#)
3. Click `File` tab -> Click `Save a copy in drive` button
 - This will save the notebook file in your google drive, which is necessary to follow the tutorial
 - Check `Colab Notebooks` directory in your [google drive](#) if the notebook is saved successfully

Common steps for training a neural network in Colab

1. Connect to your google drive
2. Import modules
3. Configure the experiments (e.g. hyper-parameters)
4. Construct data pipeline
5. Construct a neural network builder
6. Initialize the network and optimizer
7. Load pre-trained weight if exists
8. Train the network
9. Visualize and analyze the results

1. Connect to your google drive

- This step is **required** if you want to save checkpoints into your drive and load them later on

```
from google.colab import drive  
drive.mount('/gdrive')  
gdrive_root = '/gdrive/My Drive'
```

Go to this URL in a browser: <https://accounts.google.com/o/oauth2/auth?c>

Enter your authorization code:

.....

Mounted at /gdrive

2. Import modules

```
import os

import torch
import torch.optim as optim
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import MNIST
```

3. Configure the experiments

```
# training & optimization hyper-parameters
max_epoch = 10
learning_rate = 0.0001
batch_size = 200
device = 'cuda'

# model hyper-parameters
input_dim = 784 # 28x28=784
hidden_dim = 512
output_dim = 10
```


4. Construct data pipeline

- **torchvision.datasets.MNIST** will automatically construct **MNIST** dataset.
- **torch.utils.data.DataLoader** receives MNIST dataset and does followings
 - parse data using multi-processing
 - make mini-batches of data
 - shuffle data when make a mini-batch

```
data_dir = os.path.join(gdrive_root, 'my_data')  
  
transform = transforms.ToTensor()  
  
train_dataset = MNIST(data_dir, train=True, download=True, transform=transform)  
train_dataloader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True, drop_last=True)  
  
test_dataset = MNIST(data_dir, train=False, download=True, transform=transform)  
test_dataloader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False, drop_last=False)
```

5. Construct a neural network builder

- Here we're going to train a simple MLP-based neural network with ReLU non-linear activation functions.

```
class MyClassifier(nn.Module):
    def __init__(self, input_dim=784, hidden_dim=512, output_dim=10):
        super(MyClassifier, self).__init__()
        self.layers = nn.Sequential(
            nn.Linear(input_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, hidden_dim),
            nn.ReLU(),
            nn.Linear(hidden_dim, output_dim),
        )

    def forward(self, x):
        batch_size = x.size(0)
        x = x.view(batch_size, -1)
        outputs = self.layers(x)
        return outputs
```

6. Initialize the network and optimizer

- Initialize the network and pass its parameters to the optimizer

```
my_classifier = MyClassifier(input_dim, hidden_dim, output_dim)
my_classifier = my_classifier.to(device)

optimizer = optim.Adam(my_classifier.parameters(), lr=learning_rate)
```

7. Load pre-trained weight if exists

- Later when we train a neural network, we're going to save checkpoints periodically into the location 'gdrive/My Drive/checkpoints'.
- And if you have a saved checkpoint there, this block will load it and resume the training.

```
ckpt_dir = os.path.join(gdrive_root, 'checkpoints')
if not os.path.exists(ckpt_dir):
    os.makedirs(ckpt_dir)

ckpt_path = os.path.join(ckpt_dir, 'lastest.pt')
if os.path.exists(ckpt_path):
    ckpt = torch.load(ckpt_path)
    best_acc = ckpt['best_acc']
    my_classifier.load_state_dict(ckpt['my_classifier'])
    optimizer.load_state_dict(ckpt['optimizer'])
    print('checkpoint is loaded !')
    print('current best accuracy : %.2f' % best_acc)
else:
    best_acc = 0
```

8. Now train the network

- Training session consists of mainly three parts:
 - train phase
 - test phase
 - backup phase

```
it = 0
train_losses = []
test_losses = []
for epoch in range(max_epoch):
    # train phase
    # ...

    # test phase
    # ...

    # save checkpoint whenever there is improvement in performance
    # ...
```

8. Now train the network - train phase

```
it = 0
train_losses = []
test_losses = []
for epoch in range(max_epoch):
    # train phase
    for inputs, labels in train_dataloader:
        it += 1

        # load data to the GPU.
        inputs = inputs.to(device)
        labels = labels.to(device)

        # feed data into the network and get outputs.
        logits = my_classifier(inputs)

        # calculate loss
        # Note: `F.cross_entropy` function receives logits, or pre-softmax outputs, rather than final probability scores.
        loss = F.cross_entropy(logits, labels)

        # Note: You should flush out gradients computed at the previous step before computing gradients at the current step.
        # Otherwise, gradients will accumulate.
        optimizer.zero_grad()

        # backpropagate loss.
        loss.backward()

        # update the weights in the network.
        optimizer.step()

        # calculate accuracy.
        acc = (logits.argmax(dim=1) == labels).float().mean()

        if it % 200 == 0:
            print('[epoch:{}, iteration:{}] train loss : {:.4f} train accuracy : {:.4f}'.format(epoch, it, loss.item(), acc.item()))

    # save losses in a list so that we can visualize them later.
    train_losses.append(loss.item())
```

8. Now train the network - test phase

- test phase follows the same steps as the train phase, except that:
 - use test data instead of train data
 - do not back-propagate loss and update weights

```
# test phase
n = 0.
test_loss = 0.
test_acc = 0.
for test_inputs, test_labels in test_dataloader:
    test_inputs = test_inputs.to(device)
    test_labels = test_labels.to(device)

    logits = my_classifier(test_inputs)
    test_loss += F.cross_entropy(logits, test_labels, reduction='sum')
    test_acc += (logits.argmax(dim=1) == test_labels).float().sum()
    n += inputs.size(0)

test_loss /= n
test_acc /= n
test_losses.append(test_loss.item())
print('[epoch:{}, iteration:{}] test_loss : {:.4f} test accuracy : {:.4f}'.format(epoch, it, test_loss.item(), test_acc.item()))
```

8. Now train the network - checkpointing

- It is always a good idea to save the checkpoints periodically, otherwise you'll lose everything you've trained if the session is expired.

```
# save checkpoint whenever there is improvement in performance
if test_acc > best_acc:
    best_acc = test_acc
    # Note: optimizer also has states ! don't forget to save them as well.
    ckpt = {'my_classifier':my_classifier.state_dict(),
            'optimizer':optimizer.state_dict(),
            'best_acc':best_acc}
    torch.save(ckpt, ckpt_path)
    print('checkpoint is saved !')
```


8. Now train the network - results

- After 10 epochs,
you'll get near 98% accuracy on MNIST dataset


```
[epoch:5, iteration:1600] train loss : 0.1125 train accuracy : 0.9700
[epoch:5, iteration:1800] train loss : 0.1012 train accuracy : 0.9700
[epoch:5, iteration:1800] test_loss : 0.1256 test accuracy : 0.9606
checkpoint is saved !
[epoch:6, iteration:2000] train loss : 0.0906 train accuracy : 0.9750
[epoch:6, iteration:2100] test_loss : 0.1069 test accuracy : 0.9663
checkpoint is saved !
[epoch:7, iteration:2200] train loss : 0.0742 train accuracy : 0.9750
[epoch:7, iteration:2400] train loss : 0.0898 train accuracy : 0.9750
[epoch:7, iteration:2400] test_loss : 0.1042 test accuracy : 0.9671
checkpoint is saved !
[epoch:8, iteration:2600] train loss : 0.0897 train accuracy : 0.9800
[epoch:8, iteration:2700] test_loss : 0.0972 test accuracy : 0.9698
checkpoint is saved !
[epoch:9, iteration:2800] train loss : 0.0885 train accuracy : 0.9700
[epoch:9, iteration:3000] train loss : 0.0652 train accuracy : 0.9750
[epoch:9, iteration:3000] test_loss : 0.0889 test accuracy : 0.9720
checkpoint is saved !
```

9. Visualize and analyze the results - 1

```
import matplotlib.pyplot as plt
```

```
plt.plot(train_losses, label='train loss')  
plt.plot(test_losses, label='test loss')  
plt.legend()
```

<matplotlib.legend.Legend at 0x7f0183010be0>



9. Visualize and analyze the results - 2

```
import random
from PIL import Image

num_test_samples = len(test_dataset)
random_idx = random.randint(0, num_test_samples)

topil = transforms.transforms.ToPILImage()
test_input, test_label = test_dataset.__getitem__(random_idx)
test_prediction = F.softmax(my_classifier(test_input.unsqueeze(0).to(device)), dim=1).argmax().item()
print('label : %i' % test_label)
print('prediction : %i' % test_prediction)

test_image = topil(test_input)
test_image.resize((128, 128))
```

label : 0
prediction : 0

